

Lecture 4: The GP example

Verena Zuber

15/2/2019

```
library(lme4)
library(lattice)
library(ggplot2)
library(nlme)
```

```
data.chol = read.csv("Cholesterol.txt", sep=" ")
dim(data.chol)
```

```
## [1] 441 6
```

```
head(data.chol)
```

```
## chol doctor age bmi agedoc sex
## 1 7.13 1 54 27.39 55 0
## 2 7.70 1 55 29.10 55 0
## 3 7.30 1 56 27.90 55 0
## 4 6.89 1 71 26.67 55 1
## 5 6.90 1 72 26.70 55 1
## 6 7.90 1 73 29.70 55 1
```


1. First part: Fitting the model

Model 1 (individual level)

```
model1 = lm(chol ~ age, data=data.chol)
summary(model1)
```

```
##
## Call:
## lm(formula = chol ~ age, data = data.chol)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.8971 -0.6206 -0.1105  0.5693  2.9456
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  2.798691   0.268571   10.42  <2e-16 ***
## age          0.051262   0.004301   11.92  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.8362 on 439 degrees of freedom
## Multiple R-squared:  0.2445, Adjusted R-squared:  0.2428
## F-statistic: 142.1 on 1 and 439 DF, p-value: < 2.2e-16


data.chol$PooledPredictions = fitted(model1)
plot(data.chol$age, data.chol$chol, xlab = "age", ylab = "total cholesterol", cex = 1.2)
abline(model1, col = "red", lwd = 2)
```



Model 2 (group level)

```
chol.group = tapply(data.chol$chol, INDEX=data.chol$doctor, FUN=mean)
age.group = tapply(data.chol$age, INDEX=data.chol$doctor, FUN=mean)
Group.Model = lm(chol.group ~ age.group)
summary(Group.Model)
```

```
##
## Call:
## lm(formula = chol.group ~ age.group)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.7216 -0.4513 -0.1844  0.3020  1.3576
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   1.30687    5.05233   0.259   0.801
## age.group     0.07548    0.08176   0.923   0.378
##
## Residual standard error: 0.67 on 10 degrees of freedom
## Multiple R-squared:  0.07854,    Adjusted R-squared:  -0.0136
## F-statistic: 0.8524 on 1 and 10 DF,  p-value: 0.3776
plot(age.group, chol.group, xlab = "age", ylab = "total cholesterol", cex = 1.2)
abline(Group.Model, col = "red", lwd = 2)
```




Model 3 (individual regression, one regression model for each group)

```
fits = lmList(chol ~ age | doctor, data=data.chol)
fits
```

```
## Call:
##   lmerMod::lmList(data = data.chol, chol ~ age | doctor)
##   Data: data.chol
##
## Coefficients:
##   (Intercept)          age
## 1      4.3314316  0.04143871
## 2      5.4010239  0.03006138
## 3      3.5442585  0.04202075
## 4      0.7656213  0.08916495
## 5      3.0816670  0.04737634
## 6      2.9456065  0.05329581
## 7      0.1915603  0.08862860
## 8      4.6621327  0.01243695
## 9      2.7119802  0.04830460
## 10     1.5642149  0.06244402
## 11     2.5339182  0.04815023
## 12     2.7996834  0.03780979
##
## Degrees of freedom: 441 total; 417 residual
## Residual standard error: 0.5597412
```

```
xyplot(chol ~ age | doctor, data=data.chol, type = c("p", "r"),
       col.line = "darkorange", xlab = "age", ylab = "total cholesterol")
```



Model 4 (fixed effects)

```
FE.Model = lm(chol ~ age + as.factor(doctor), data=data.chol)
summary(FE.Model)
```

```
##
## Call:
## lm(formula = chol ~ age + as.factor(doctor), data = data.chol)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.59881 -0.40321 -0.08463  0.37929  1.77313
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    3.826236   0.213854  17.892 < 2e-16 ***
## age            0.049543   0.003065  16.164 < 2e-16 ***
## as.factor(doctor)2  0.400993   0.136014   2.948  0.00337 **
## as.factor(doctor)3 -0.752146   0.135865  -5.536 5.41e-08 ***
## as.factor(doctor)4 -0.555317   0.133254  -4.167 3.73e-05 ***
```

```
## as.factor(doctor)5 -0.884528 0.136039 -6.502 2.21e-10 ***
## as.factor(doctor)6 -0.653299 0.135970 -4.805 2.15e-06 ***
## as.factor(doctor)7 -1.295580 0.133444 -9.709 < 2e-16 ***
## as.factor(doctor)8 -1.563657 0.136053 -11.493 < 2e-16 ***
## as.factor(doctor)9 -1.193645 0.135970 -8.779 < 2e-16 ***
## as.factor(doctor)10 -1.453255 0.133231 -10.908 < 2e-16 ***
## as.factor(doctor)11 -1.376027 0.136039 -10.115 < 2e-16 ***
## as.factor(doctor)12 -1.685593 0.137173 -12.288 < 2e-16 ***
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.5764 on 428 degrees of freedom
## Multiple R-squared: 0.65, Adjusted R-squared: 0.6402
## F-statistic: 66.24 on 12 and 428 DF, p-value: < 2.2e-16
```

Model 5 (random intercept)

```
RandomIntercept = lme( chol ~ age, random = ~ 1 | doctor, data = data.chol )
summary(RandomIntercept)
```

```
## Linear mixed-effects model fit by REML
## Data: data.chol
##      AIC      BIC    logLik
## 828.697 845.035 -410.3485
##
## Random effects:
## Formula: ~1 | doctor
##      (Intercept) Residual
## StdDev: 0.6347908 0.5764246
##
## Fixed effects: chol ~ age
##              Value Std.Error DF t-value p-value
## (Intercept) 2.9060357 0.26477408 428 10.97553      0
## age          0.0495831 0.00306279 428 16.18888      0
## Correlation:
##      (Intr)
## age -0.714
##
## Standardized Within-Group Residuals:
##      Min      Q1      Med      Q3      Max
## -2.7850636 -0.7012544 -0.1419474 0.6536599 3.0850909
##
## Number of Observations: 441
## Number of Groups: 12
```

Model 6 (random intercept and slope)

```
RandomSlope = lme( chol ~ age, random = ~ 1+age | doctor, data = data.chol )
summary(RandomSlope)
```

```
## Linear mixed-effects model fit by REML
## Data: data.chol
##      AIC      BIC    logLik
## 821.9886 846.4956 -404.9943
##
## Random effects:
```

```
## Formula: ~1 + age | doctor
## Structure: General positive-definite, Log-Cholesky parametrization
##           StdDev      Corr
## (Intercept) 1.28163791 (Intr)
## age         0.01771585 -0.872
## Residual    0.55997509
##
## Fixed effects: chol ~ age
##           Value Std.Error DF t-value p-value
## (Intercept) 2.8791744 0.4215200 428 6.830458      0
## age         0.0500704 0.0060597 428 8.262837      0
## Correlation:
## (Intr)
## age -0.901
##
## Standardized Within-Group Residuals:
##           Min      Q1      Med      Q3      Max
## -2.8523390 -0.6664557 -0.1141926  0.6206844  3.0809630
##
## Number of Observations: 441
## Number of Groups: 12
```

2. Second part: Using ggplot

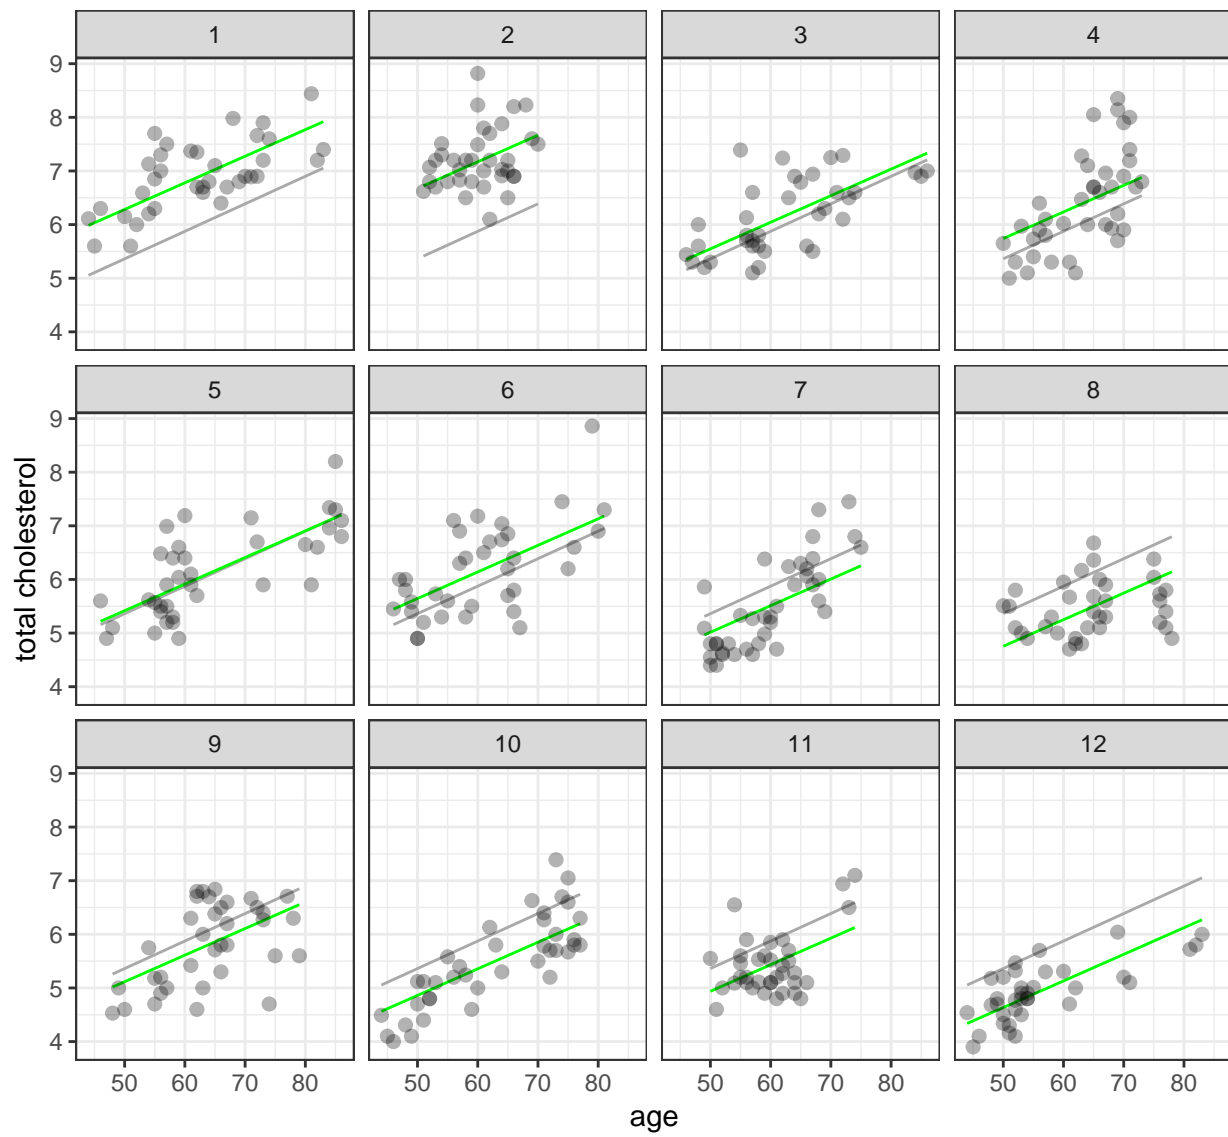
In order to use ggplot we need to add the model fits (`fitted()`) from the random effects (Model 5-7) to the data sets.

```
data.chol$FixedEffects = fitted(RandomIntercept)
data.chol$VaryingInterceptPredictions = fitted(RandomIntercept)
data.chol$VaryingSlopePredictions = fitted(RandomSlope)
```

Pooled prediction and fixed effects

```
gg.fe <- ggplot(data.chol, aes(x = age, y = chol, group = doctor)) +
  geom_line(aes(y = PooledPredictions), color = "darkgrey") +
  geom_line(aes(y = FixedEffects), color = "green") +
  geom_point(alpha = 0.3, size = 2) +
  xlab("age") + ylab("total cholesterol") +
  facet_wrap(~doctor) +
  theme_bw()

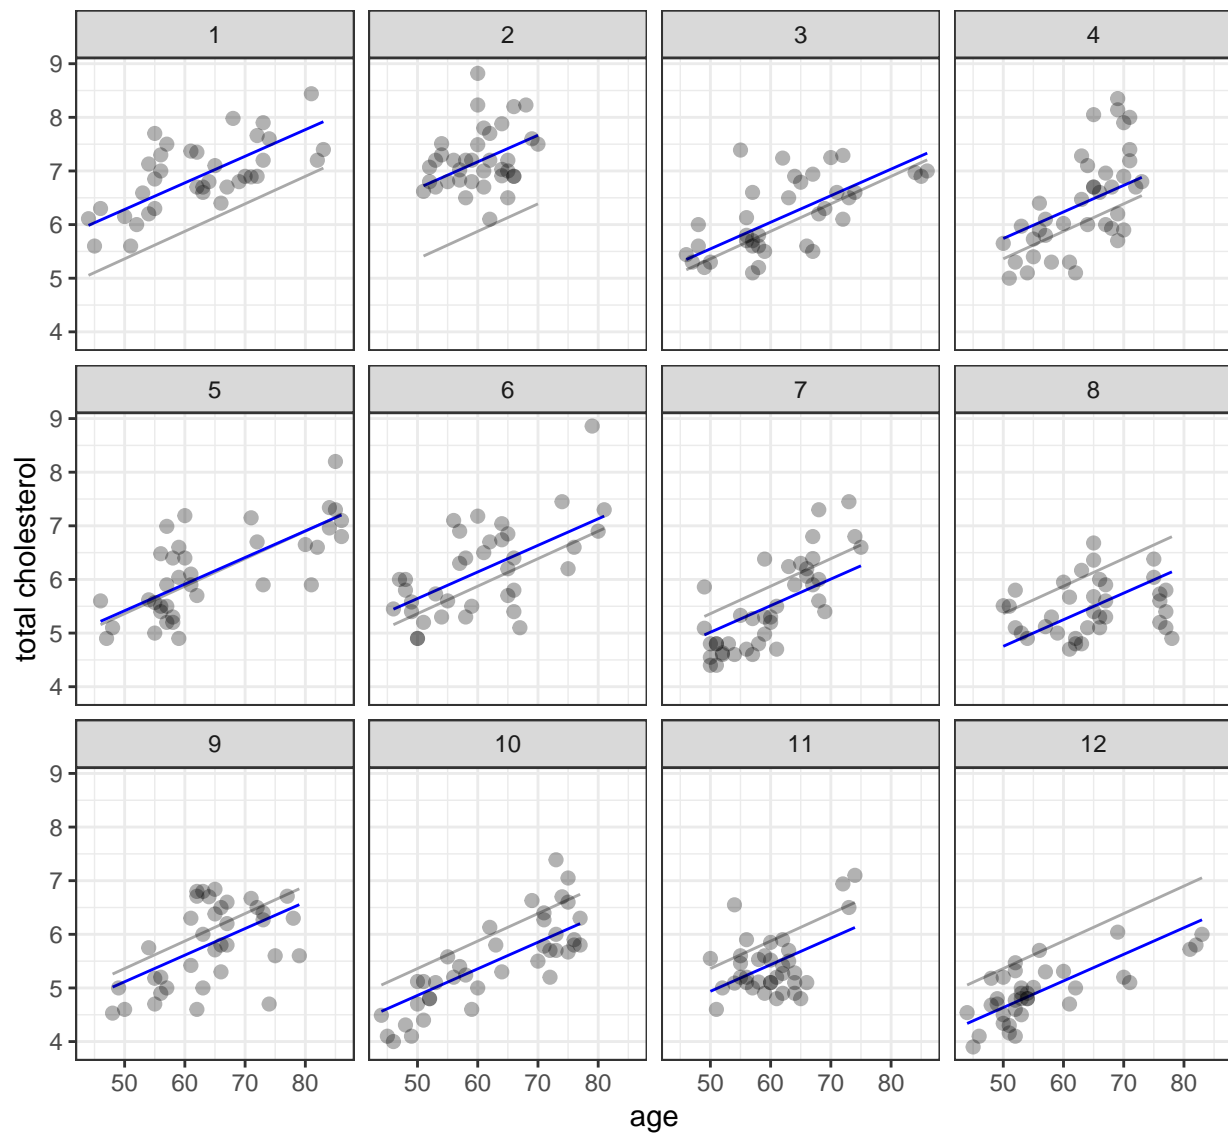
print(gg.fe)
```



Pooled prediction and random intercept

```
gg.intercept <- ggplot(data.chol, aes(x = age, y = chol, group = doctor)) +
  geom_line(aes(y = PooledPredictions), color = "darkgrey") +
  geom_line(aes(y = VaryingInterceptPredictions), color = "blue") +
  geom_point(alpha = 0.3, size = 2) +
  xlab("age") + ylab("total cholesterol") +
  facet_wrap(~doctor) +
  theme_bw()

print(gg.intercept)
```



Pooled prediction and random intercept and slope

```
gg.slope <- ggplot(data.chol, aes(x = age, y = chol, group = doctor)) +
  geom_line(aes(y = PooledPredictions), color = "darkgrey") +
  geom_line(aes(y = VaryingSlopePredictions), color = "red") +
  geom_point(alpha = 0.3, size = 2) +
  xlab("age") + ylab("total cholesterol") +
  facet_wrap(~doctor) +
  theme_bw()

print(gg.slope)
```