

# Introduction to Deep Learning

Elizaveta Semenova

Imperial College London  
19 March 2024

# Outline

## Motivation

A neuron

## Neural networks

Network design

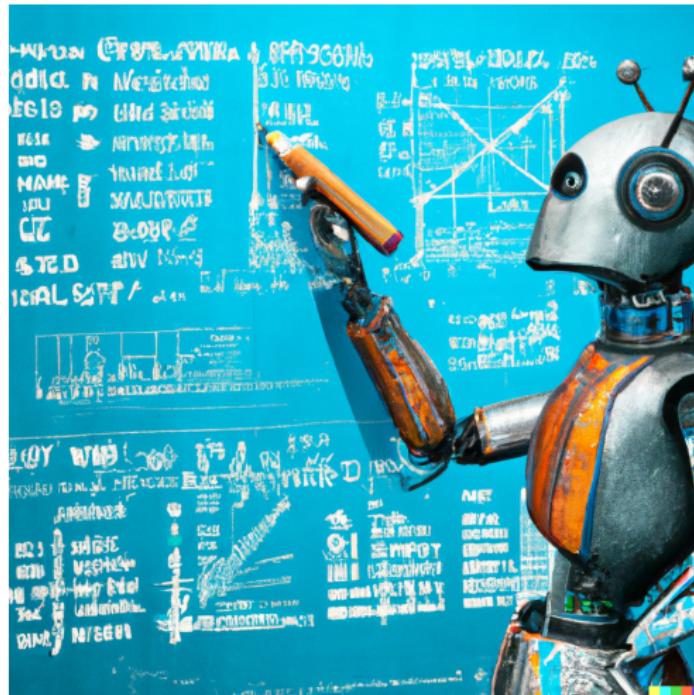
Activation functions

Training a neural network

## Deep learning tools

## A practical example

# Modern applications of machine learning



"Robot mixing statistics, calculus and linear algebra, digital art", DALLE-2

# Modern applications of machine learning

ChatGPT 3.5 ▾



You

Why are neural networks useful? In one sentence



ChatGPT

Neural networks are useful because they can learn complex patterns and relationships from data, enabling tasks such as image recognition, natural language processing, and predictive modeling.



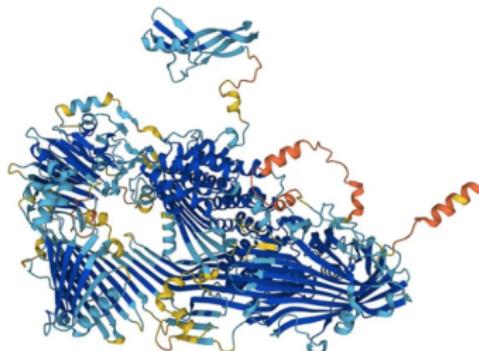
# Modern applications of machine learning

NEWS | 18 January 2024

## AlphaFold found thousands of possible psychedelics. Will its predictions help drug discovery?

Researchers have doubted how useful the AI protein-structure tool will be in discovering medicines – now they are learning how to deploy it effectively.

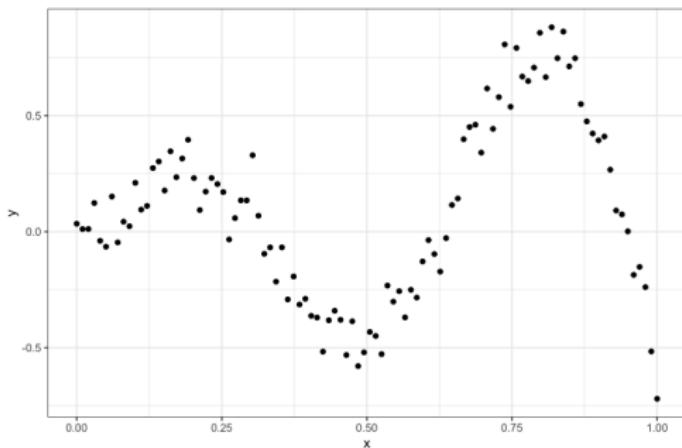
By [Ewen Callaway](#)



Protein structures predicted by AlphaFold have helped to identify candidate drug compounds. Credit: DeepMind

# Curve fitting

- ▶ Problem setting: given **observed** data pairs  $(x_i, y_i), i = 1, \dots, n$ , learn to predict  $y$  from  $x$ .



# Curve fitting

- ▶ Problem setting: given **observed** data pairs  $(x_i, y_i), i = 1, \dots, n$ , learn to predict  $y$  from  $x$ .
- ▶ Model  $f$ , depending on **parameters**  $\theta$ :  $y = f(x, \theta)$ .

# Curve fitting

- ▶ Problem setting: given **observed** data pairs  $(x_i, y_i), i = 1, \dots, n$ , learn to predict  $y$  from  $x$ .
- ▶ Model  $f$ , depending on **parameters**  $\theta$ :  $y = f(x, \theta)$ .

linear model	$f(x, \theta) = \theta^T x$	parametric, linear
Gaussian process	$f(x, \theta) = GP_\theta(x)$	nonparametric
neural network	$f(x, \theta) = NN_\theta(x, \theta)$	parametric, non-linear

# Curve fitting

- ▶ Problem setting: given **observed** data pairs  $(x_i, y_i), i = 1, \dots, n$ , learn to predict  $y$  from  $x$ .
- ▶ Model  $f$ , depending on **parameters**  $\theta$ :  $y = f(x, \theta)$ .
- ▶ **Training** a model, means finding parameters  $\theta^*$ , such that  $f(x_i, \theta^*) \approx y_i$ , e.g.:

$$\sum_{i=1}^n (f(x_i, \theta^*) - y_i)^2 \rightarrow \min_{\theta}$$

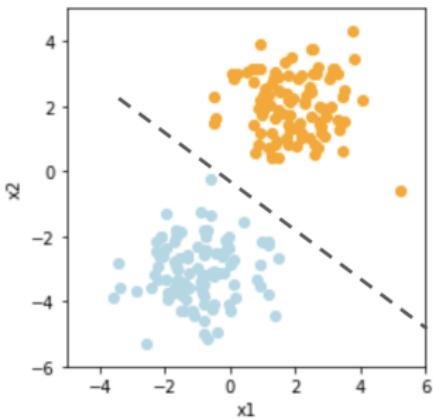
# Curve fitting

$$\underbrace{\sum_{i=1}^n (f(x_i, \theta^*) - y_i)^2}_{\mathcal{L}(x, y, \theta)} \rightarrow \min_{\theta}$$

$\mathcal{L}(x, y, \theta)$  - loss function

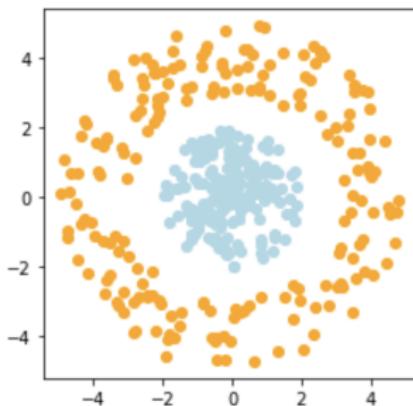
# Classification problem

- ▶ Logistic regression is a **classification** model.
- ▶ It allows to separate two classes using a **linear** decision boundary.



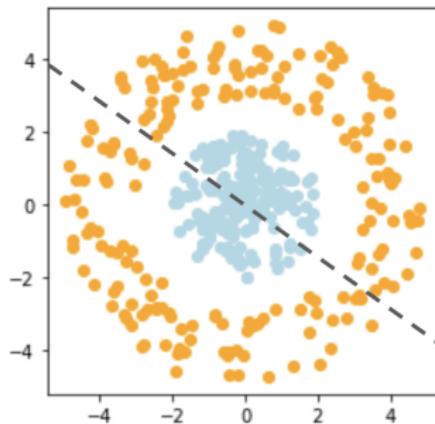
## Classification problem

- ▶ What would happen if we apply logistic regression to this data?



## Classification problem

- ▶ What would happen if we apply logistic regression to this data?
- ▶ Not possible, the task is too complex!



# Outline

Motivation

A neuron

Neural networks

Network design

Activation functions

Training a neural network

Deep learning tools

A practical example

# Neurons in a brain

- ▶ The brain uses billions of interconnected neurons.
- ▶ Many simple interactions can solve complex tasks.
- ▶ Can we use this concept to make computers learn?

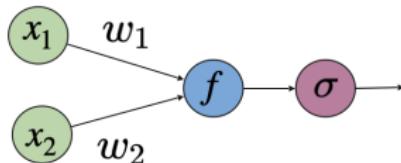


# Computational model of a single neuron

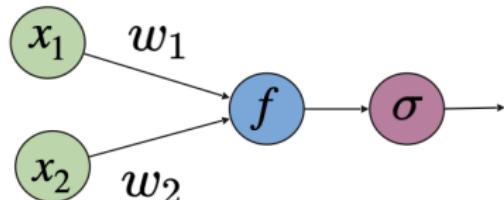
- ▶ Inputs  $x_1, x_2$  are multiplied with **weights**  $w_1, w_2$  and added together, and a **bias**  $b$  is added:

$$f(x_1, x_2) = x_1 * w_1 + x_2 * w_2 + 1 * b$$

- ▶ We then apply an **activation function**.



# Computational model of a single neuron



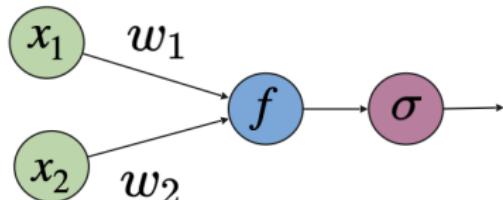
$$f(x_1, x_2) = x_1 w_1 + x_2 w_2 + b,$$

$$p(x_1, x_2) = \sigma(f(x_1, x_2)),$$

$$\sigma(t) = \frac{1}{1 + \exp^{-t}}$$

- ▶ Does it remind you of anything?

# Computational model of a single neuron



$$f(x_1, x_2) = x_1 w_1 + x_2 w_2 + b,$$

$$p(x_1, x_2) = \sigma(f(x_1, x_2)),$$

$$\sigma(t) = \frac{1}{1 + \exp^{-t}}$$

- ▶ Does it remind you of anything?
- ▶ Logistic regression!

# Outline

Motivation

A neuron

**Neural networks**

Network design

Activation functions

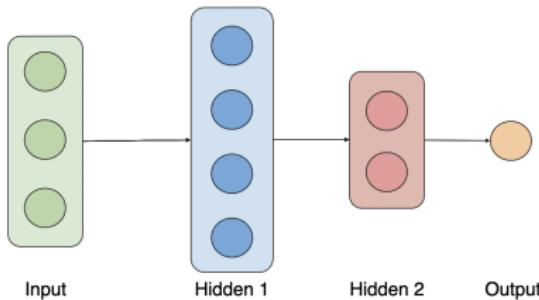
Training a neural network

Deep learning tools

A practical example

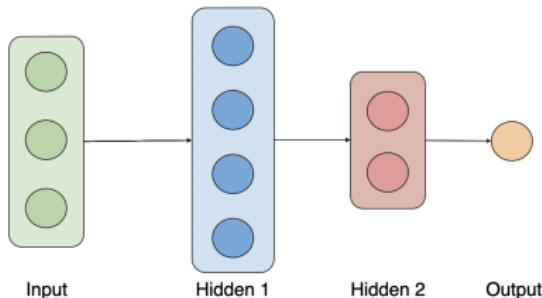
# A neural network

- ▶ A neural network connects many individual neurons (**nodes**) together to solve complex tasks.
- ▶ The **architecture** of a neural network tells us how these connections should look like.



# A neural network

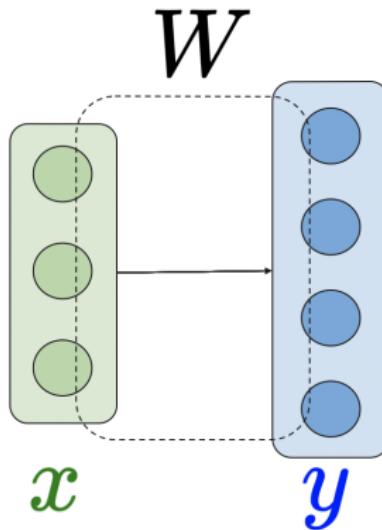
- ▶ The **input layer** consists of data - predictors, also called features.
- ▶ The **output layer** is the observed variable of interest.
- ▶ In-between are the **hidden layers**.



# Mathematical foundations of deep learning

- ▶ How to derive an appropriate loss function for each problem?  
→ **Statistics**
- ▶ How to optimise a loss function?  
→ **Calculus**
- ▶ How to express computations efficiently?  
→ **Linear algebra**

# Neural networks in matrix notation

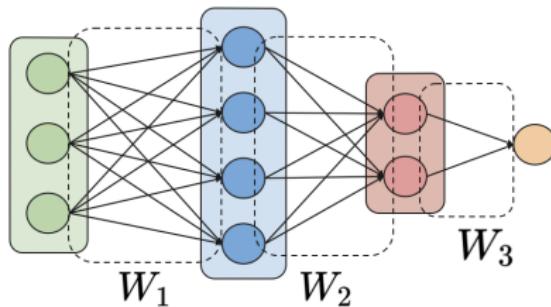


$$y = Wx$$

# Network design

We refer to the design of the network as the **network architecture**:

- ▶ The number of hidden layers,
- ▶ The number of neurons per layer (referred to as width),
- ▶ The activation function for each neuron.

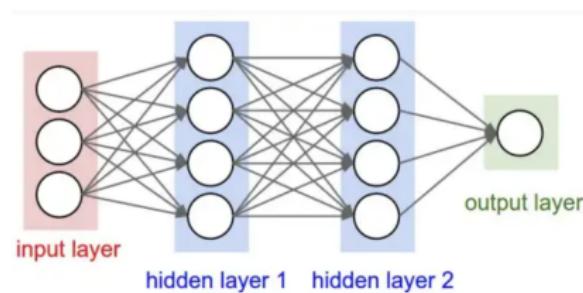


# Architectures

The list of useful architectures is evolving daily as they are being discovered by trial-and-error.

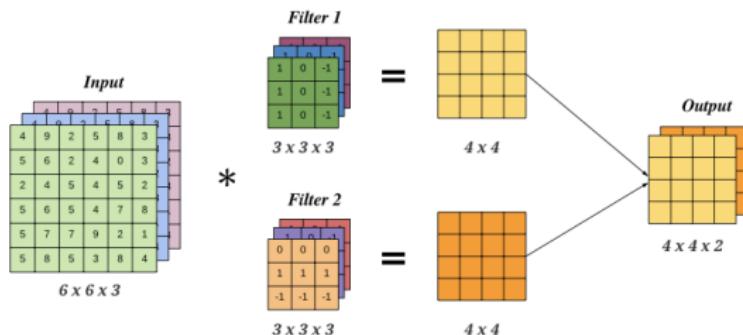
# Multilayer perceptron (MLP)

- ▶ A **multilayer perceptron** is a feedforward architecture, consisting of fully connected neurons with a nonlinear kind of activation function.



# Convolutional neural networks (CNNs)

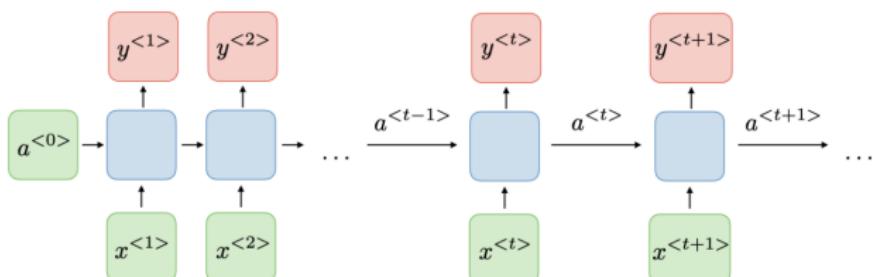
- ▶ CNNs use **convolution layers** which scan the input with a set of **filters**.
- ▶ **Pooling layers** are used for downsampling.



Credit: Kevin Murphy, "Probabilistic Machine Learning 2"

# Recurrent neural networks (RNNs)

- ▶ RNNs allow previous outputs to be used as inputs while having hidden states. They are used to model **sequences**.



Credit: Shervine Amidi, Stanford CS230 course

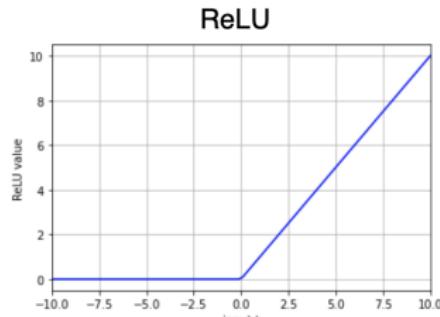
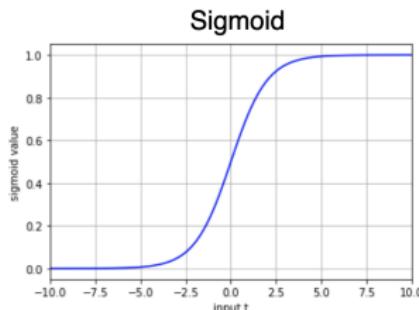
# Activation functions

- ▶ Every node in the network has an **activation function**.
- ▶ To perform binary classification, our last node has a **sigmoid** activation

$$\sigma(t) = \frac{1}{1 + \exp -t}$$

- ▶ For nodes in the hidden layers we usually apply the Rectified Linear Unit (**ReLU**)

$$\text{ReLU}(t) = \max(0, t)$$



# Training a neural network

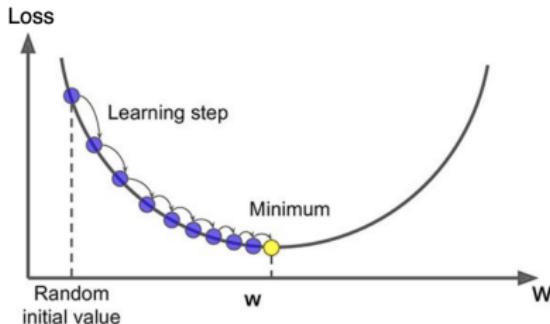
- ▶ **To train** a neural network means to find weights  $W$  that **minimise** the loss function  $\mathcal{L}(W)$ .

# Gradient descent

- ▶ **Gradient descent** can be used to iteratively update the weights and get the best loss value

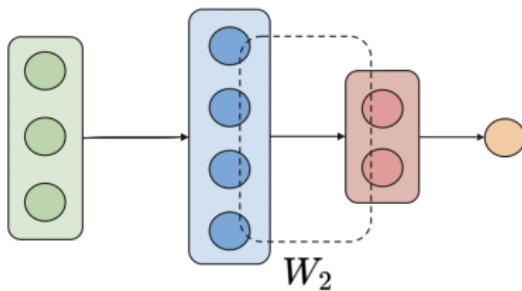
$$w_i := w_i - \alpha \frac{\partial \mathcal{L}}{\partial w_i}$$

- ▶ The step size is determined by the **learning rate**  $\alpha$ .



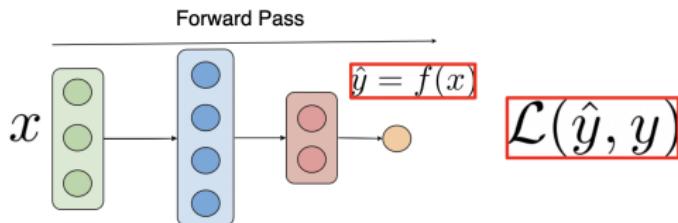
# Back propagation

How can we update the hidden layers based on the loss value determined from the last layer?



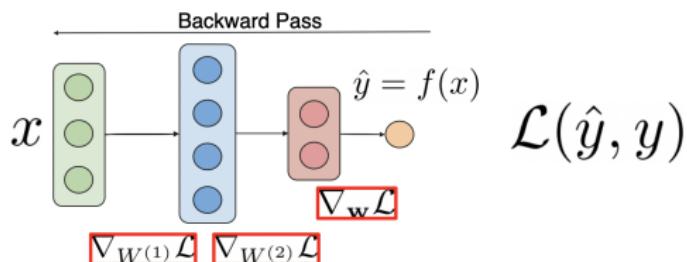
## Forward pass

- ▶ Propagate inputs  $x$  through the network to compute the **predicted value**  $\hat{y}$ .
- ▶ Compute the **loss value**  $\mathcal{L}(y, \hat{y})$  between the prediction  $\hat{y}$  and the known ground-truth value  $y$ .



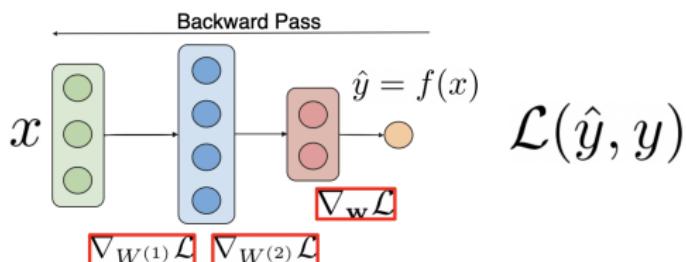
# Backward pass

- ▶ Given the loss, compute the derivatives with respect to each layer ("back-propagation").
- ▶ Update the weights of each layer based on the computed gradients.



# The training loop

- ▶ Define architecture.
- ▶ Load data.
- ▶ For a given number of iterations:
  - ▶ Compute model prediction  $\hat{y}$ ,
  - ▶ Compute loss value  $\mathcal{L}(y, \hat{y})$ ,
  - ▶ Perform back-propagation to obtain the gradients of the loss,
  - ▶ Update the networks parameters based on the gradient.



# Outline

Motivation

A neuron

Neural networks

Network design

Activation functions

Training a neural network

Deep learning tools

A practical example

# Deep learning tools: Python



## Deep learning libraries: R

In R, there are some native libraries, as well as interfaces to libraries written in other languages:

- ▶ *neuralnet*,
- ▶ *deepnet*,
- ▶ *h2o*,
- ▶ *KerasR*

and more.

For examples, see

<https://srdas.github.io/DLBook/DeepLearningWithR.html>

# Outline

Motivation

A neuron

Neural networks

Network design

Activation functions

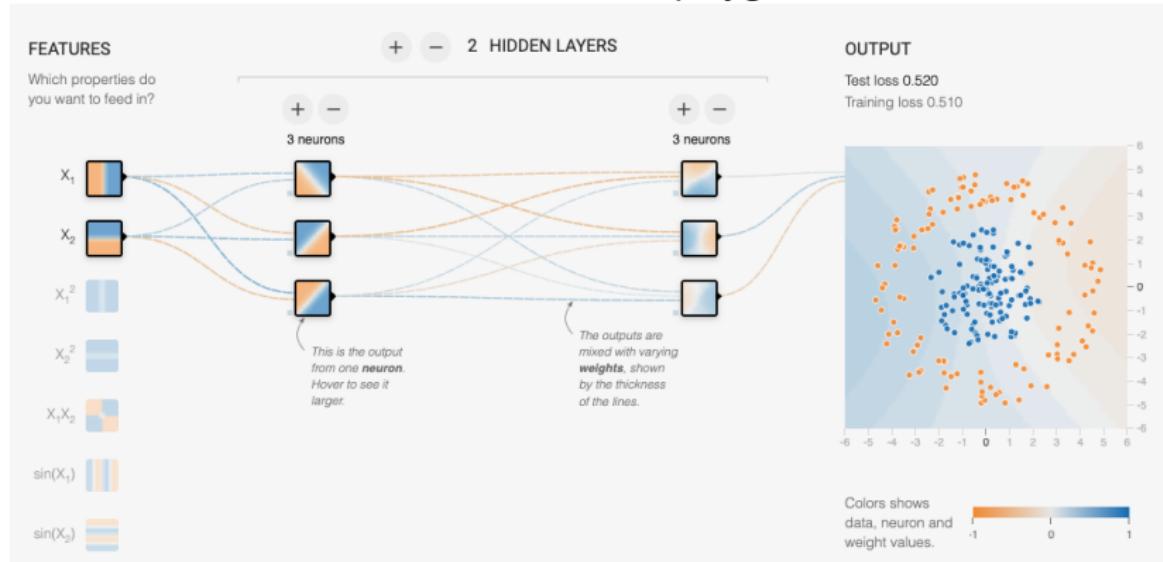
Training a neural network

Deep learning tools

A practical example

# Try for yourself

Search for 'tensorflow playground':



## Example: binary classification

```
# ----- simulate data -----

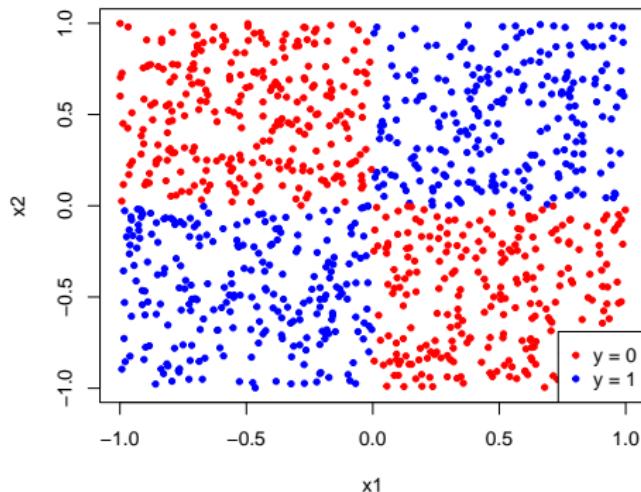
# Generate some sample data
set.seed(123)
n <- 1000
x1 <- runif(n, min = -1, max = 1)
x2 <- runif(n, min = -1, max = 1)
y <- ifelse((x1 * x2) > 0, 1, 0) # binary classification task

# Visualize dependence between X and y
plot(x1, x2, col = ifelse(y == 1, "blue", "red"), pch = 20,
      main = "Dependence between X and y",
      main = "",
      xlab = "x1", ylab = "x2")
legend("bottomright", legend = c("y = 0", "y = 1"), col = c("red", "blue"), pch = 20)

# Combine into a data frame
data <- data.frame(x1, x2, y)
```

## Example: binary classification

Let's fit a binary classification model to this simulated dataset:



## Example: define neural network

```
# ----- define neural network -----
# Install and load required packages
library(neuralnet)

# Split data into training and testing sets
train_indices <- sample(1:n, 0.7 * n)
train_data <- data[train_indices, ]
test_data <- data[-train_indices, ]

# Define the neural network architecture
# Here, we have 2 input nodes (x1 and x2), 2 nodes in the hidden layer, and 1 output node
# We use 'sigmoid' as the activation function for both the hidden and output layers
nn <- neuralnet(y ~ x1 + x2, data = train_data, hidden = 5, act.fct = "logistic")
```

# Example: evaluate neural network

```
# ----- evaluate neural network -----

# Predict on test data
predicted <- predict(nn, test_data)

# Convert predicted probabilities to binary predictions
predicted_class <- ifelse(predicted > 0.5, 1, 0)

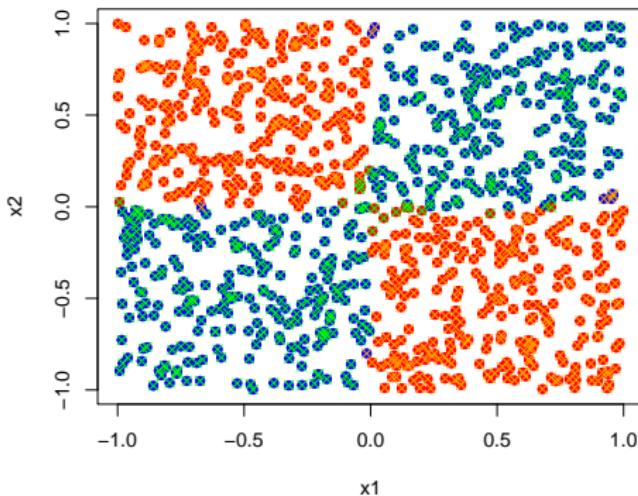
# Calculate accuracy
accuracy <- mean(predicted_class == test_data$y)
cat("Accuracy:", accuracy, "\n")

predicted_full <- predict(nn, data)
y_pred_full <- ifelse(predicted_full > 0.5, 1, 0)

# Visualize dependence between X and y
plot(x1, x2, col = ifelse(y == 1, "blue", "red"), pch = 19,
      main = "Dependence between X and y",
      xlab = "x1", ylab = "x2")
points(x1, x2, col = ifelse(y_pred_full == 1, "green", "orange"), pch = 4)
```

## Example: evaluate neural network

How well does the model perform?



## Example: code

This code is available at

[https://github.com/elizavetasemenova/nن\\_in\\_R](https://github.com/elizavetasemenova/nن_in_R)

## How to learn more?

- ▶ "Dive into deep learning", Zhang, Smola, Lipton, Li (2023)  
<https://d2l.ai/>
- ▶ "Understanding Deep Learning", Prince (2023)
- ▶ Stanford, CS230 course  
<https://stanford.edu/~shervine/teaching/cs-230/>
- ▶ "Deep learning with R", <https://srdas.github.io/DLBook/DeepLearningWithR.html>

# Thank You!